



Integer Factorization as an Ising problem

Anil Prabhakar, Advith Desu, Aryan Namboodri

Indian Institute of Technology Madras

May 11, 2026

Factorization Problem

- ▶ Given a natural number N , find P, Q such that

$$N = PQ$$

- ▶ Hard inverse problem underlying RSA security.
- ▶ Map factorization $\rightarrow$ QUBO (or Ising) problem.
- ▶ Solve using annealing-based algorithms.

Binary Representation and Assumptions

Assumptions:

- ▶ N is odd.
- ▶ Factors have equal bit-length

$$n = \left\lceil \frac{1}{2} \log_2 N \right\rceil.$$

Binary form:

$$P = \sum_{k=0}^{n-1} 2^k p_k, \quad Q = \sum_{k=0}^{n-1} 2^k q_k, \quad p_k, q_k \in \{0, 1\}.$$

Fixed bits:

$$p_0 = q_0 = 1, \quad p_{n-1} = q_{n-1} = 1.$$

Column-wise Multiplication

- ▶ Construct binary long multiplication table.
- ▶ Introduce carry variables $s_{i,j}$.
- ▶ Each column gives a constraint:

$$C_i = N_i - \sum_j q_j p_{i-j} - \sum s_{j,i} + \sum 2^j s_{i,j}.$$

- ▶ Correct factors satisfy $C_i = 0$.

Optimization objective:

$$\min_{\{p_i, q_i, s_{ij}\}} \sum_i (C_i)^2.$$

- ▶ Binary variables: factor bits + carry bits.
- ▶ Ground state corresponds to valid factorization.
- ▶ Alternative direct method:

$$\left(N - \sum p_i 2^i \sum q_j 2^j\right)^2$$

but gives fewer preprocessing opportunities.

Rule 1: Saturated Sum

For binary $x_i \in \{0, 1\}$:

$$\sum_{i=1}^n x_i = n \implies x_i = 1 \forall i.$$

Why: the only way n binary variables sum to n is for every one of them to be 1.

A symmetric form with negative coefficients:

$$\sum_{i=1}^n -x_i = -n \implies x_i = 1 \forall i.$$

Rule 2: Zero Sum

For binary $x_i \in \{0, 1\}$:

$$\sum_{i=1}^n x_i = 0 \implies x_i = 0 \forall i.$$

Why: a sum of non-negative binaries equals 0 only when every term is 0.

The same conclusion holds with negative coefficients:

$$\sum_{i=1}^n -x_i = 0 \implies x_i = 0 \forall i.$$

Rule 3: Doubling Rule

For binary variables $x_1, x_2, x_3 \in \{0, 1\}$:

$$x_1 + x_2 = 2x_3 \implies x_1 = x_2 = x_3.$$

Why: $\text{LHS} \in \{0, 1, 2\}$ and $\text{RHS} \in \{0, 2\}$. Equality forces LHS even, so $x_1 + x_2 \in \{0, 2\}$ — both equal, and equal to x_3 .

Use in factorization: appears whenever a column clause has only two factor-bit terms balanced against a single carry, e.g.

$$p_k + q_k = 2s \Rightarrow p_k = q_k = s.$$

Rule 4: Coefficient Bound Rule

General linear constraint on $x_i \in \{0, 1\}$:

$$\sum_i c_i x_i + C = 0, \quad S_+ = \sum_{c_i > 0} c_i, \quad S_- = \sum_{c_j < 0} |c_j|.$$

If a single coefficient is larger than the max possible opposing sum, the variable is *forced*:

Positive coefficients ($c_i > 0$):

$$c_i > S_- - C \Rightarrow x_i = 0, \quad c_i > S_+ + C \Rightarrow x_i = 1.$$

Negative coefficients ($c_j < 0$, $|c_j|$ magnitude):

$$|c_j| > S_+ + C \Rightarrow x_j = 0, \quad |c_j| > S_- - C \Rightarrow x_j = 1.$$

Example: $x_1 + x_2 = 2x_3 + 1$. Coefficient of x_3 is 2, max LHS is 2, so x_3 cannot be 1 $\Rightarrow x_3 = 0$.

Rule 5: Complement Rule

For binary $x_1, x_2 \in \{0, 1\}$:

$$x_1 + x_2 = 1 \implies x_2 = 1 - x_1, \quad x_1 x_2 = 0.$$

Why: the only solutions are $(0, 1)$ and $(1, 0)$ — exactly one variable is 1, so their product is always 0.

Use in factorization: eliminates one variable outright (substitute $x_2 \rightarrow 1 - x_1$) and kills any quadratic term $x_1 x_2$ that appears in a column clause.

Parity Rule

Both sides of a clause must have the same parity:

$$x_1 + x_2 + \text{even} = \text{even} \Rightarrow x_1 = x_2,$$

$$x_1 + x_2 + \text{even} = \text{odd} \Rightarrow x_1 + x_2 = 1.$$

Example:

$$x_1 + x_2 + 2x_3 = 2y_1 + 4y_2 \Rightarrow x_1 = x_2,$$

$$x_1 + x_2 + 2x_3 = 2y_1 + 4y_2 + 1 \Rightarrow x_1 + x_2 = 1.$$

Why it matters: column clauses balance a small sum of factor bits against carries with even coefficients ($2s_{i,i+1}, 4s_{i,i+2}, \dots$). The parity of the factor-bit sum is therefore locked by N_i .

Power Rule

For any binary variable $x \in \{0, 1\}$:

$$x^m = x, \quad m = 2, 3, \dots$$

Why: $0^m = 0$ and $1^m = 1$ for all positive integers m .

Use in factorization: after squaring the energy function $E = \sum_i C_i^2$, every p_k^2 , q_k^2 , s^2 collapses back to its linear form p_k , q_k , s . This is what makes the squared objective expressible as a (degree-2) QUBO once cubics/quartics are quadrized.

Worked Example: $N = 391 = 17 \times 23$

Setup:

- ▶ $391 = 110000111_2$
- ▶ **Fixed bits:** $p_0 = q_0 = 1$ (odd), $p_4 = q_4 = 1$ (MSB)
- ▶ **Free bits:** p_1, p_2, p_3 and q_1, q_2, q_3 — 6 unknowns

Long multiplication table (cross-product $p_a q_b$ sits in column $a + b$):

	2^8	2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0
P					1	p_3	p_2	p_1	1
Q					1	q_3	q_2	q_1	1
				q_1	$p_3 q_1$	$p_2 q_1$	$p_1 q_1$	q_1	1
			q_2	$p_3 q_2$	$p_2 q_2$	$p_1 q_2$	q_2		
		q_3	$p_3 q_3$	$p_2 q_3$	$p_1 q_3$	q_3			
	1	p_3	p_2	p_1	1				
Carry-ins	$s_{6,8}, s_{7,8}$	$s_{5,7}, s_{6,7}$	$s_{4,6}, s_{5,6}$	$s_{3,5}, s_{4,5}$	$s_{2,4}, s_{3,4}$	$s_{2,3}$	$s_{1,2}$		
Carry-outs		$s_{7,8}$	$s_{6,7}, s_{6,8}$	$s_{5,6}, s_{5,7}$	$s_{4,5}, s_{4,6}$	$s_{3,4}, s_{3,5}$	$s_{2,3}, s_{2,4}$	$s_{1,2}$	
$N = 391$	1	1	0	0	0	0	1	1	1

Carry $s_{i,j}$ = bit of column i 's carry destined for column j (coefficient 2^{j-i}). In each C_k , positive-sign carries enter column k (**carry-ins**); negative-sign carries with weights $2, 4, \dots$ leave column k (**carry-outs**).

Column Clauses: $C_i = 0$

Each column i : (column sum) + (carries in) = $N_i + 2s_{i,i+1} + 4s_{i,i+2}$.

$$C_0 : 1 - 1 = 0 \quad \checkmark$$

$$C_1 : p_1 + q_1 - 1 - 2s_{1,2} = 0$$

$$C_2 : p_2 + q_2 + p_1q_1 + s_{1,2} - 1 - 2s_{2,3} - 4s_{2,4} = 0$$

$$C_3 : p_3 + q_3 + p_1q_2 + p_2q_1 + s_{2,3} - 2s_{3,4} - 4s_{3,5} = 0$$

$$C_4 : 2 + p_1q_3 + p_3q_1 + p_2q_2 + s_{2,4} + s_{3,4} - 2s_{4,5} - 4s_{4,6} = 0$$

$$C_5 : p_1 + q_1 + p_2q_3 + p_3q_2 + s_{3,5} + s_{4,5} - 2s_{5,6} - 4s_{5,7} = 0$$

$$C_6 : p_2 + q_2 + p_3q_3 + s_{4,6} + s_{5,6} - 2s_{6,7} - 4s_{6,8} = 0$$

$$C_7 : p_3 + q_3 + s_{5,7} + s_{6,7} - 1 - 2s_{7,8} = 0$$

$$C_8 : 1 + s_{6,8} + s_{7,8} - 1 = 0$$

Applying Reduction Rules

C_8 (**Rule 2, sum = 0**): $s_{6,8} + s_{7,8} = 0 \Rightarrow \boxed{s_{6,8} = s_{7,8} = 0}$.

C_2 (**Rule 4, large coefficient**): LHS terms $p_2 + q_2 + p_1 q_1 + s_{1,2} \leq 4$, but coefficient of $s_{2,4}$ is 4. If $s_{2,4}=1$, need LHS ≥ 5 . $\boxed{s_{2,4} = 0}$.

C_1 (**Rule 5**): $p_1 + q_1 = 1 + 2s_{1,2}$. Max sum = $2 < 3$, so $s_{1,2} = 0$:

$$\boxed{s_{1,2} = 0, \quad q_1 = 1 - p_1} \Rightarrow p_1 q_1 = p_1(1 - p_1) = 0$$

C_2 (**after substitutions**): $p_2 + q_2 = 1 + 2s_{2,3}$. Same logic as C_1 :

$$\boxed{s_{2,3} = 0, \quad q_2 = 1 - p_2}$$

C_7 (**Rule 2, sum = 1 exactly-one**):

$$p_3 + q_3 + s_{5,7} + s_{6,7} = 1 \quad (s_{7,8} = 0 \text{ already})$$

Exactly one of these four binaries equals 1.

Remaining Structure After Pre-processing

After the rules above, we have eliminated 5 variables ($s_{1,2}, s_{2,3}, s_{2,4}, s_{6,8}, s_{7,8}$) and reduced q_1, q_2 to functions of p_1, p_2 .

Free QUBO variables (13):

$\{p_1, p_2, p_3, q_3\} \cup \{s_{3,4}, s_{3,5}, s_{4,5}, s_{4,6}, s_{5,6}, s_{5,7}, s_{6,7}\}$ (plus 2 auxiliaries from quadrization)

Remaining clauses (after substituting $q_1=1-p_1, q_2=1-p_2$):

$$\begin{aligned} C_3 : p_3 + q_3 + p_1(1-p_2) + p_2(1-p_1) - 2s_{3,4} - 4s_{3,5} &= 0 \\ &= p_3 + q_3 + p_1 + p_2 - 2p_1p_2 - 2s_{3,4} - 4s_{3,5} = 0 \end{aligned}$$

$$\begin{aligned} C_4 : 2 + p_1q_3 + p_3(1-p_1) + p_2(1-p_2) + s_{3,4} - 2s_{4,5} - 4s_{4,6} &= 0 \\ &= 2 + p_1q_3 + p_3 - p_1p_3 + s_{3,4} - 2s_{4,5} - 4s_{4,6} = 0 \end{aligned}$$

$$C_5 : 1 + p_2q_3 + p_3(1-p_2) + s_{3,5} + s_{4,5} - 2s_{5,6} - 4s_{5,7} = 0$$

$$C_6 : 1 + p_3q_3 + s_{4,6} + s_{5,6} - 2s_{6,7} = 0$$

$$C_7 : p_3 + q_3 + s_{5,7} + s_{6,7} = 1$$

Cubic terms appear after squaring (e.g. $p_1p_2 \cdot s_{3,4}$ in C_3^2 , $p_1p_3 \cdot s_{4,5}$ in C_4^2) $\Rightarrow$ quadrization required.

Quadrization and QUBO Construction

Energy function: $E = \sum_{i=3}^7 C_i^2$. Squaring introduces cubics like $a \cdot p_i p_j s_{k,\ell}$.

Rosenberg substitution: replace each AB pair appearing in a cubic by an auxiliary w , with penalty

$$P(A, B, w) = M(AB - 2Aw - 2Bw + 3w), \quad M \gg \max |\text{coeff}|.$$

$P = 0$ iff $w = AB$. After substitution every cubic $AB \cdot s$ becomes $w \cdot s$ (quadratic).

Auxiliaries needed: $w_1 = p_1 p_2$, $w_2 = p_1 p_3$.

Full QUBO:

$$H = \sum_{i=3}^7 C_i^2 \Big|_{w_1 \rightarrow p_1 p_2, w_2 \rightarrow p_1 p_3} + M[P(p_1, p_2, w_1) + P(p_1, p_3, w_2)]$$

Active QUBO variables (13 total):

Original (4)	p_1, p_2, p_3, q_3
Carry (7)	$s_{3,4}, s_{3,5}, s_{4,5}, s_{4,6}, s_{5,6}, s_{5,7}, s_{6,7}$
Auxiliary (2)	$w_1 = p_1 p_2, w_2 = p_1 p_3$

QUBO Matrix Structure

Sketch of nonzero entries of Q for $N = 391$ (after pre-processing and quadrization):

	p_1	p_2	p_3	q_3	$s_{3,4}$	$s_{3,5}$	$s_{4,5}$	$s_{4,6}$	$s_{5,6}$	$s_{5,7}$	$s_{6,7}$	w_1	w_2
p_1	•	•	•	•								•	•
p_2		•	•	•	•	•						•	
p_3			•	•	•		•	•					•
q_3				•			•	•	•	•	•		
$s_{3,4}$					•	•	•	•					
$s_{3,5}$						•	•		•	•			
$s_{4,5}$							•		•	•			
$s_{4,6}$								•	•		•		
$s_{5,6}$									•		•	•	
$s_{5,7}$										•	•	•	
$s_{6,7}$											•	•	
w_1												•	
w_2													•

The ground state of this QUBO recovers $P=17$, $Q=23$ with $E=0$.

QUBO $\rightarrow$ Ising Conversion

QUBO uses $x_i \in \{0, 1\}$; Ising-form annealers (D-Wave, our GPU SA) use spins $\sigma_i \in \{-1, +1\}$.

Linear change of variables:

$$\sigma_i = 1 - 2x_i \quad \Longleftrightarrow \quad x_i = \frac{1 - \sigma_i}{2}.$$

Substitute into the QUBO $H = \sum_i Q_{ii} x_i + \sum_{i < j} Q_{ij} x_i x_j$ and collect terms:

$$H(\sigma) = \underbrace{\sum_i h_i \sigma_i}_{\text{linear field}} + \underbrace{\sum_{i < j} J_{ij} \sigma_i \sigma_j}_{\text{couplings}} + \text{const.}$$

Coefficients (treating Q symmetrically so $Q_{ij} = Q_{ji}$ for $i \neq j$):

$$\boxed{J_{ij} = \frac{1}{4} Q_{ij}}, \quad \boxed{h_i = -\frac{1}{2} Q_{ii} - \frac{1}{4} \sum_{j \neq i} Q_{ij}}.$$

J matrix Visualization




Figure: Sparsity of J matrix for 162 spins (Integer 39203).

Sparse Coupling Matrix Representation (CSR)

Problem: The Ising coupling matrix J_{ij} is very large but mostly zero. Storing it as a dense $N \times N$ matrix wastes memory and bandwidth.

Key idea: Store *only the nonzero couplings* using **Compressed Sparse Row (CSR)** format. CSR represents a sparse matrix using three arrays:

`row_ptr[i]` start index of row i

`col_idx[k]` column index of nonzero entry

`values[k]` value of J_{ij}

Only nonzero entries are stored.

Memory cost: $\mathcal{O}(N + \text{nnz})$ instead of $\mathcal{O}(N^2)$ (nnz = number of nonzeros).

CSR Example

Consider a symmetric 5×5 coupling matrix with 8 non-zero entries (out of 25):

$$J = \begin{pmatrix} \cdot & 2 & \cdot & \cdot & -1 \\ 2 & \cdot & \cdot & 3 & \cdot \\ \cdot & \cdot & \cdot & \cdot & 4 \\ \cdot & 3 & \cdot & \cdot & \cdot \\ -1 & \cdot & 4 & \cdot & \cdot \end{pmatrix}$$

($\cdot$ denotes a stored zero — not represented in CSR)

CSR representation:

index	0	1	2	3	4	5	6	7
col_idx	1	4	0	3	4	1	0	2
values	2	-1	2	3	4	3	-1	4

index	0	1	2	3	4	5
row_ptr[i]	0	2	4	5	6	8

Reading row i : entries in `row_ptr[i]` represent the start of the i^{th} row in the `J` values array. Entries are stored at indices `row_ptr[i] ... row_ptr[i+1]-1` of the column index.

Memory cost: $\mathcal{O}(N + \text{nnz})$ instead of $\mathcal{O}(N^2)$ (nnz = number of nonzeros).

GPU Kernel




Figure: GPU Threads and Thread blocks.

GPU Thread Calculation in DA1

Each thread i calculates the local field $L_i = \sum_j J_{ij}s_j + h_i$ and the energy difference is then calculated as $-2 \times s_i \times L_i$ depending on which i^{th} spin flipped.




Figure: Energy calculation in DA1.

Graph-Coloring Simulated Annealing

In one Metropolis sweep we want to flip many spins in parallel, but two spins coupled by a non-zero J_{ij} share an edge and cannot be flipped simultaneously without invalidating each other's ΔE .

We use graph coloring on the interaction graph (nodes = spins, edges = nonzero J_{ij}) partitions spins into colors such that *no two same-color spins are coupled*. All spins in one color can then be flipped in parallel without affecting each other's energy calculations.

Two performance bins based on row degree $d_i = |\{j : J_{ij} \neq 0\}|$:

- ▶ **Dense** ($d_i \geq 128$): full block of 1024 threads, shared-memory reduction.
- ▶ **Sparse** ($d_i < 128$): one warp (32 threads) per spin, 32 spins packed per block, warp-shuffle reduction.

Number of Total and Densely Connected Spins with Scale




Figure: Number of bits in N vs number of densely connected spins.

Algorithm Pre-processing steps before annealing

Require: CSR of symmetric J , linear field h , num_spins N

- 1: Verify J symmetry and a zero diagonal
 - 2: Bin each spin: **dense** if $d_i \geq 128$, else **sparse**
 - 3: $\text{color}[] \leftarrow$ Welsh–Powell greedy coloring (largest-degree-first, first-fit)
 - 4: Bucket spins by color: $\text{color_dense_offsets}$, $\text{color_sparse_offsets}$
 - 5: Upload all CSR + color tables + hub arrays to GPU
 - 6: Initialize a random spin configuration and compute initial energy
 - 7: Calculate the start-temperature: $T_0 = \langle \Delta E^+ \rangle / \ln(1/\chi)$
(χ = target acceptance rate)
-


Annealing Loop (GPU)

Algorithm Color-parallel simulated annealing


Require: β -schedule $\{\beta_t\}$ (geometric), sweeps per β , N_{colors}

```
1: for  $\beta$  in  $\beta$ -schedule do
2:   for sweep = 1 ...  $m$  do
3:     Generate fresh uniform randoms for all  $N$  spins
4:     for  $c = 0 \dots N_{\text{colors}} - 1$  do ▷ colors run sequentially
5:       collectFlipCandidates(dense) on color  $c$ 
6:       collectFlipCandidates(sparse) on color  $c$ 
7:       each spin  $i$  in color  $c$  in parallel:
8:          $L_i \leftarrow \sum_j J_{ij}s_j + h_i$ ,  $\Delta E_i \leftarrow -2s_i L_i$ 
9:         if  $\text{rand}_i < \min(1, e^{-\beta \Delta E_i})$ : append  $(i, \Delta E_i)$  to candidates
10:      applyAllFlipsInColor: flip every accepted spin, atomicAdd  $\Delta E_i$ 
        to total energy
11:      if any hub spin in color  $c$ : refresh hub-values cache
12:    end for
13:  end for
14:  if total energy < best energy then snapshot best spins
15:  end if
16: end for
17: return best-energy spin configuration
```

Annealing Time per Iteration vs Number of Spins




Time Taken for Construction and Annealing



Each Annealing run consists of 10 sweeps at each temperature step, with the starting temperature T_0 calculated to give an initial acceptance rate of 0.5, and a geometric cooling schedule with factor 0.95 until $T < 10^{-3}$.

Solution Quality Measured by Gap between Real and Estimated P



1. Average (%) difference with square root N as the estimate = **10.33%**
2. Average (%) difference with P_{pred} from Simulated Annealing = **4.23%**
3. Average (%) difference with mean of P_{pred} and N/Q_{pred} = **2.98%**

Brute-force Search

The annealer returns approximate factors P_{pred} , Q_{pred} — usually off by a few percent. A brute-force search around this prediction finds the exact factor.

Best starting estimate — average of two independent guesses for P :

$$P_{\text{centre}} = \frac{1}{2}(P_{\text{pred}} + \lfloor N/Q_{\text{pred}} \rfloor).$$

Outward spiral: test candidates closest to P_{centre} first, alternating above and below until the exact factor is found.

Candidates are grouped into chunks of C . Even chunks spiral upward (steps up), odd chunks spiral downward (steps down). Threads claim the next chunk via an atomic counter, so the search is partially parallel.

Stop condition: the first thread to find $p \mid N$ sets a shared flag; the rest exit at the next iteration.

Brute-force Optimizations

Testing all numbers in a chunk is inefficient as we can filter out many candidates without a divisibility test using a wheel sieve cheaply. The wheel sieve eliminates candidates that are divisible by small primes.

Wheel Sieve:

- ▶ Any p divisible by 2, 3, 5, 7, or 11 cannot be the (odd, ≥ 13) prime factor.
- ▶ Using a Wheel sieve skips approximately 79% of candidates without a divisibility test.

Early termination: atomic stop flag checked at each chunk boundary; all T threads exit within one chunk of the discovery.


Brute-force Algorithm

Algorithm Parallel spiral divisor search

Require: N , P_{pred} , Q_{pred} , T threads

- 1: $P_{\text{centre}} \leftarrow \frac{1}{2}(P_{\text{pred}} + \lfloor N/Q_{\text{pred}} \rfloor)$
 - 2: **shared:** atomic chunk_idx $\leftarrow 0$, atomic stop $\leftarrow \text{false}$, atomic found $\leftarrow 0$
 - 3: **for** each thread **in parallel do**
 - 4: **while not** stop **do**
 - 5: $k \leftarrow \text{chunk_idx.fetchAdd}(1)$ ▷ claim next chunk
 - 6: **if** k even **then** ▷ step up
 - 7: $[\text{lo}, \text{hi}] \leftarrow [P_{\text{centre}} + \frac{k}{2}C, P_{\text{centre}} + (\frac{k}{2}+1)C - 1]$
 - 8: **else** ▷ step down
 - 9: $[\text{lo}, \text{hi}] \leftarrow [P_{\text{centre}} - \frac{k+1}{2}C, P_{\text{centre}} - \frac{k-1}{2}C - 1]$
 - 10: **end if**
 - 11: **for** each $p \in [\text{lo}, \text{hi}]$ passing the wheel sieve **do**
 - 12: **if** $N \bmod p = 0$ **then**
 - 13: found $\leftarrow p$; stop $\leftarrow \text{true}$; **return**
 - 14: **end if**
 - 15: **end for**
 - 16: **end while**
 - 17: **end for**
 - 18: **return** $P \leftarrow \text{found}$, $Q \leftarrow N/P$
-

Time Taken to Find Exact primes by Brute Force




$$\text{Time Taken (s)} = \frac{3.32 \times 10^{-9}}{\text{Threads}} \times \text{Gap (\%)} \times 2^n$$


End-to-End Run: ~ 100 -bit Semiprime

$$N = 416\,749\,328\,744\,899\,275\,664\,046\,210\,629$$

Annealing-guided pipeline



Naive $\sqrt{N}$ baseline



Factors found: $P = 573\,858\,364\,692\,161$, $Q = 726\,223\,323\,360\,389$ (72 threads)
The annealer's prediction places the brute-force search much closer to the true factor.

Better classical post-processing. Brute-force trial division is the simplest way to consume the annealer's guess ($P_{\text{pred}}, Q_{\text{pred}}$), but it makes almost no use of the prediction beyond using it as a starting point. Other Existing classical factorization methods may be able to exploit this kind of information far more effectively than a linear scan. Exploring such methods would be a natural next step.

The Bottleneck, from the brute-force timing model

$$\text{Time(s)} = \frac{3.32 \times 10^{-9}}{\text{threads}} \times \text{Gap(\%)} \times 2^n,$$

- ▶ **Threads:** This is the amount of 'compute' we can throw at the problem. Would be helpful, but its bounded by hardware.
- ▶ **Gap (%):** a multiplicative factor in front of 2^n . Since 2^n grows exponentially with bit-length, *shrinking the gap helps us keep up* with larger N .

The next step would be to improve the solution quality from the annealer by exploring better QUBO encodings, and any other optimizations that drives P_{pred} closer to P and chips away at the Gap(%).

References



B. Wang, F. Hu, H. Yao, and C. Wang, "Prime factorization algorithm based on parameter optimization of Ising model," *Scientific Reports*, vol. 10, no. 1, 2020. DOI: [10.1038/s41598-020-62802-5](https://doi.org/10.1038/s41598-020-62802-5).



S. Jiang, K. A. Britt, A. J. McCaskey, T. S. Humble, and S. Kais, "Quantum Annealing for Prime Factorization," *Scientific Reports*, vol. 8, no. 1, 2018. DOI: [10.1038/s41598-018-36058-z](https://doi.org/10.1038/s41598-018-36058-z).



T. Preis, P. Virnau, W. Paul, and J. J. Schneider, "GPU accelerated Monte Carlo simulation of the 2D and 3D Ising model," *Journal of Computational Physics*, vol. 228, no. 12, pp. 4468–4477, 2009. DOI: [10.1016/j.jcp.2009.03.018](https://doi.org/10.1016/j.jcp.2009.03.018).